

Sea Surface Temperature Prediction Mini-Project Report
Chris Lee, Minghan Sun, Yishan Xu, Yege Zhang
March 15, 2020

Objective (Yege)

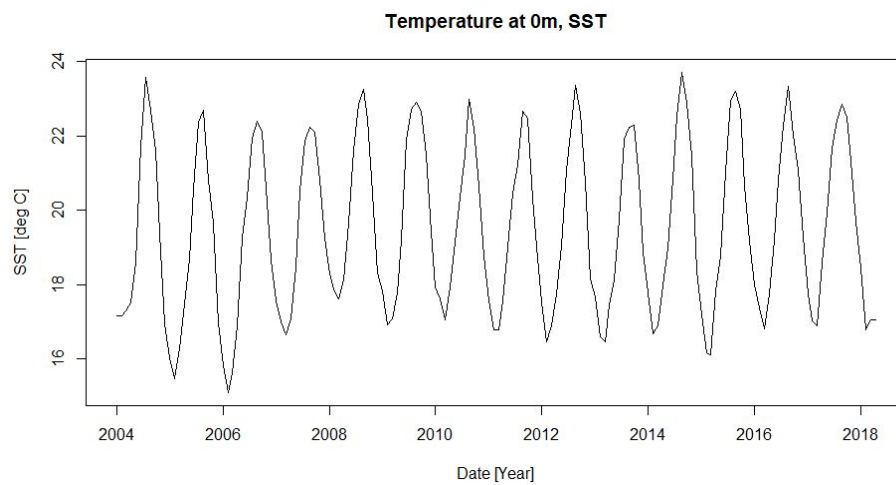
To analyze the monthly sea temperature from 0 meters to 90 meters below the sea level, in order to predict sea surface temperature in the future.

Data (Yege)

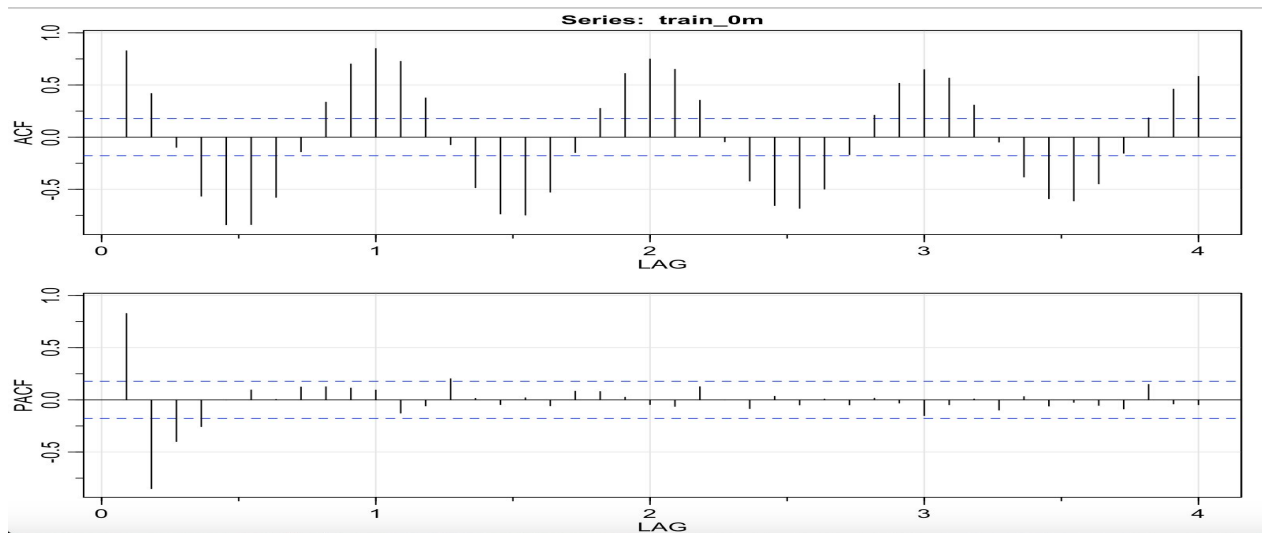
For this project, we use the monthly sea temperature from 2014 to 2018, totaling 158 observations. According to the data, it has 33 days in one period, and 11 periods in a year.

Data Exploration (Yege)

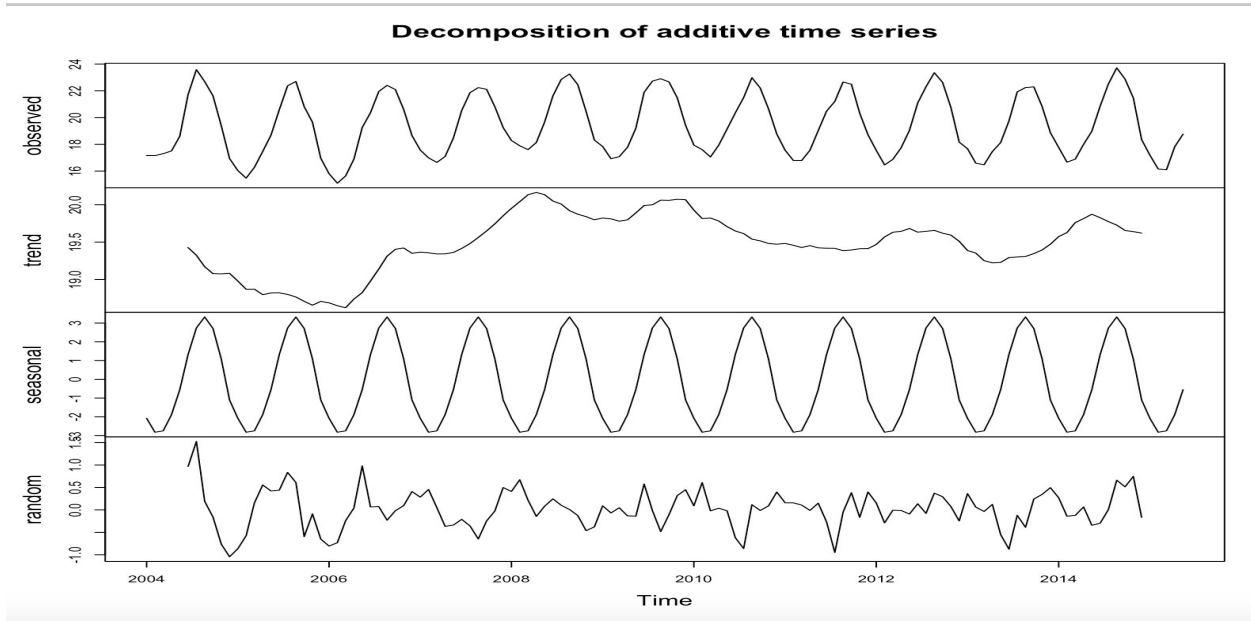
The plot below shows the original time series, which has the very obvious seasonality.




To further analyze, we plot the ACF and PACF. As we can see, the seasonality occurs in every time lag.



We then decompose the data into 3 parts: 1) trend, 2) seasonal, 3) random. As we can see in the figure, our data has the upward and downward trend and the seasonality.



In order to have stationary data, we remove the trend and seasonality from the time series. However, we are still able to see the seasonality from the plot. To try to solve the issue, we decide to try the `auto.arima` function, which will be explained in the following model.



Model 1: SARIMA without Additional Variable (Minghan)

Ideally, we would like to find an optimal SARIMA model by constantly adjusting models until no correlation is found in residuals. However, the time series of sea surface temperature is more complicated than we thought. Manually trying and adjusting multiple models based on acf/pacf plot won't give a desired result. Ultimately, we turned to 'auto.arima' function for help, where we let the machine look for the best model that fits our sea surface temperature train series.

Below is the flow of how we find the appropriate SARIMA model based on train test.

Separating data as train/test:

```
#train test split
train_0m=ts(ndata$`0m`, start = c(2004,1,1), end = c(2015,5,15), frequency = 11)
test_0m=ts(ndata$`0m`, start = c(2015,6,17), end = c(2018,4,12), frequency = 11)
```

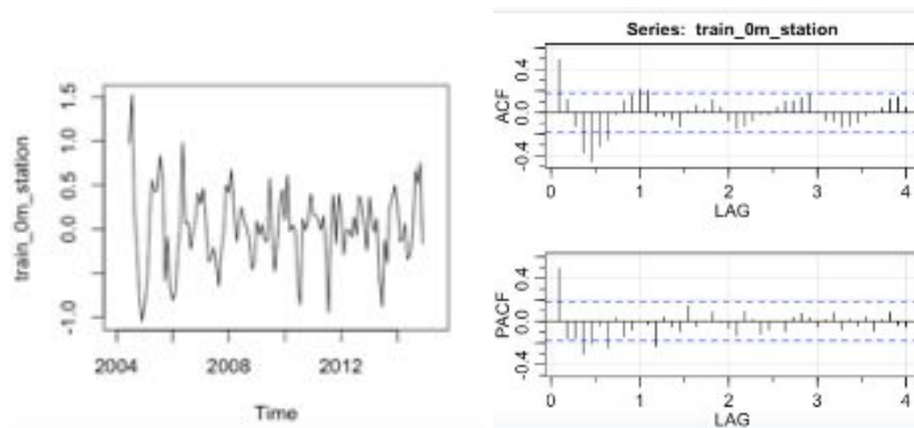
First 126 observations were put in the train test, last 32 observations were put in the test set.

Trial 1: Use adjusted series to train models.

```
decompose<-decompose(train_0m)
train_0m_station<-train_0m-decompose$seasonal-decompose$trend
```

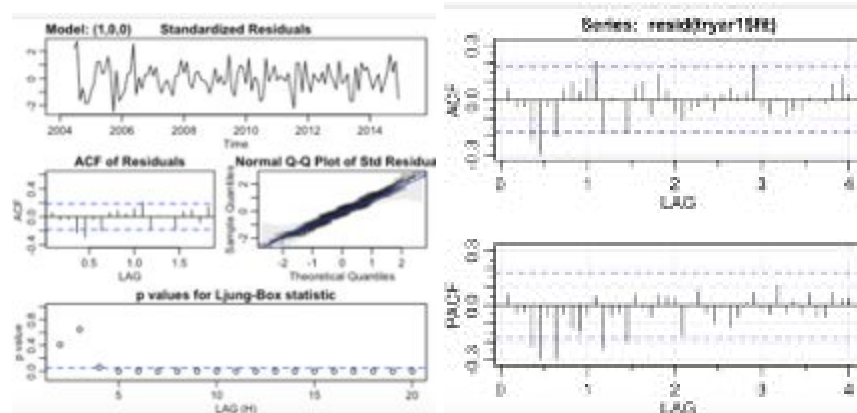
After we have taken out the trend and seasonality in the data, the series still displays uneven variance across the periods.

However, we decided to try AR1 model based on the acf/pacf graph below,

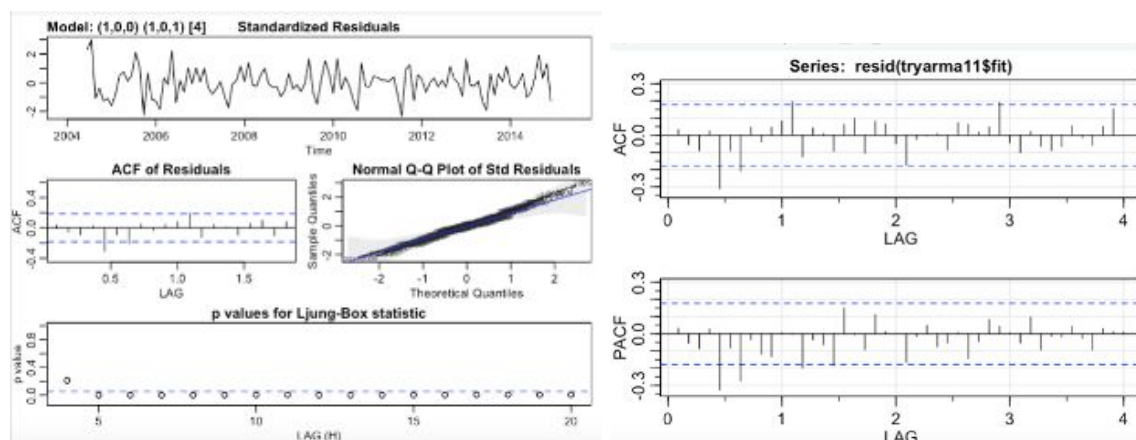


Diagnosis plot displayed below tells that there exists correlation among residuals starting from lag 4.

Checking residual plot to find the appropriate model that fits residual, we have found no clear pattern, although both ACF/PACF start to stick out at lag=4.




Adding seasonal ARMA(1,1) with lag=4 to original model. The diagnosis plot shown as below.



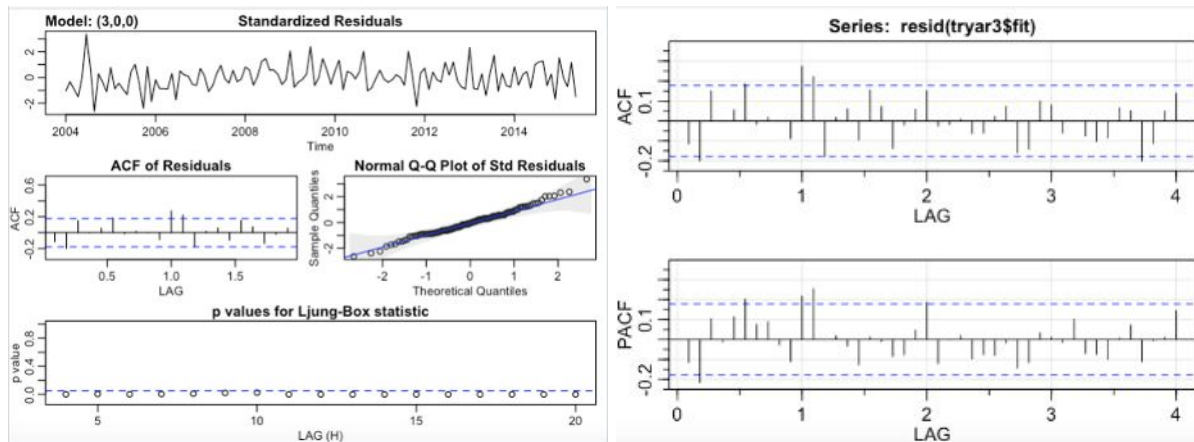
Q-statistic still shows correlation among residuals starting from lag=5. Residual plot for new model almost identical to previous model. After several 'try and error', no desired model was found.

Trial 2: Use original series to fit models

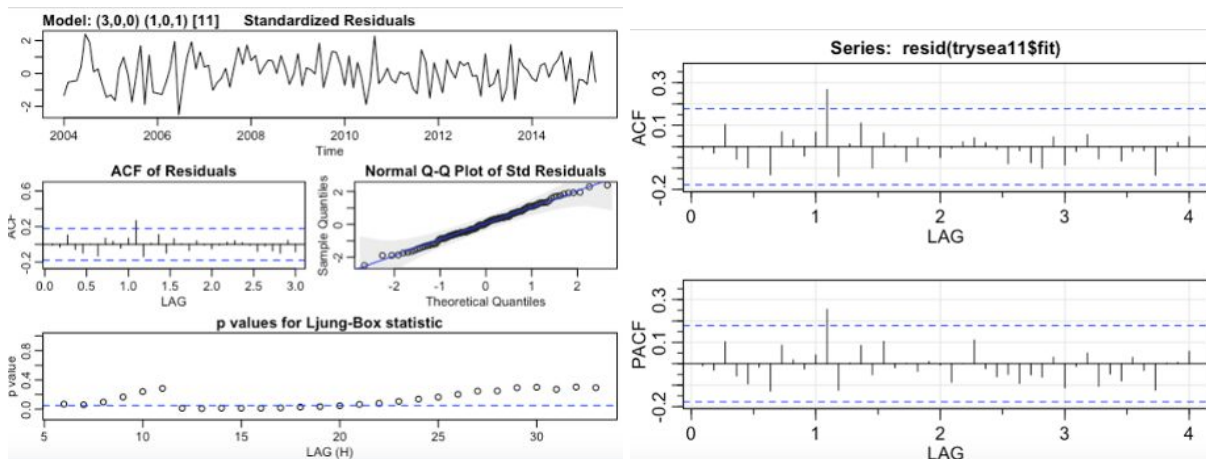
According to the ACF/PACF graph below, AR(3) seems to be a good model to try first as pacf cut off at lag=3.



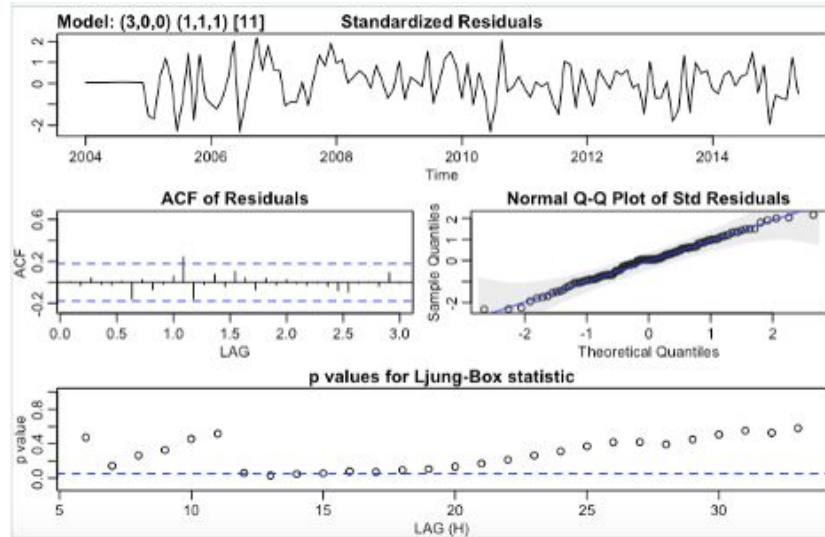
Diagnosis plot tells that there is correlation among residuals. Model needs to be adjusted.



Since ACF/PACF stick out at lag=1 (11th period), adding seasonal ARMA(1,1) with s=11 to the original AR(3) may be appropriate. New diagnosis plot and residual check are displayed below.



Clear seasonality displayed, adding seasonal difference = 1 to the model.



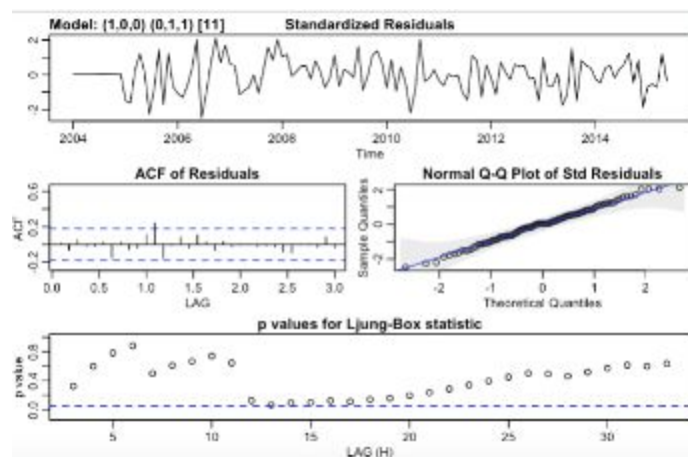
Model looks not bad, with great improvement compared to previous models. With a couple of p-value points almost below the blue line, we wanted to see if there is a better model than this. By adjusting several times, no model is better than the one above; SARIMA(3,0,0)*(1,1,1) S=11. Therefore, we decided to use 'auto.arima' function.

Trial 3: 'auto.arima' function

```
#Using auto.arima to obtain the best model
fitauto_seasonal<-auto.arima(train_0m,seasonal=TRUE,trace=TRUE,test='kpss',ic='bic')
```

Best model: ARIMA(1,0,0)(0,1,1)[11]

By allowing seasonal component, auto.arima has selected SARIMA(1,0,0)*(0,1,1) S=11 as the best model based on the train set. Diagnosis plot is shown below.



P value plot looks good and better than the last model in general. In early lags, the p-value points are further to the blue line compared to the last model. No point is below the blue line. QQ plot is almost identical to the 45 degree line.

Now, we can forecast using SARIMA(1,0,0)*(0,1,1) S=11.

Forecast 1: using 'forecast' function

When using the 'forecast' function, we allow the model to use some predicted values to predict a certain stage. The MSPE of this method is equal to 22.18.

```
> mean((forecastedvalue$mean-test_0m)^2) #mspe=22.1746  
[1] 22.1764
```

Forecast 2: using all the true value to perform one-step-ahead

In this method, we have created a function that will use all the true values prior to the predicted period to forecast. The MSPE of this method is equal to 21.99.

```
> mean((osapredictedvalue-test_0m)^2) #mspe=21.9856  
[1] 21.9856
```

Conclusion: The MSPE from two forecast methods are very close. It indicates that, in this case, whether the series we use to perform one-step-ahead contains predicted value or not does not affect much on the prediction error. On average, using SARIMA without additional variables will give us mean squared predicted error of predicted sea surface temperature around 22.

Model 2: SARIMA with additional variables (Yishan)

Read the csv file into R, rename columns as “0m”, “10m”...

```
> data <- read.csv("gilbralter_time_series_r.2.csv")
> ndata <- data[,-c(1,3)]
> names(ndata) <- c("0m", "10m", "20m", "30m", "40m", "50m", "60m", "70m", "80m", "90m")
```

Split the data into training (first 80% observations) and the testing (remaining 20% observations). Set frequency equals to 11 because one period equals 33 days, and there are 11 periods in a year

```
> train_0m=ts(ndata$`0m`, start = c(2004,1,1), end = c(2015,5,15), frequency = 11)
> test_0m=ts(ndata$`0m`, start = c(2015,6,17), end = c(2018,4,12), frequency = 11)
> |
```

As we can see from the initial model, the MSPE score is very high. Sea surface temperature in the past alone might not be able to accurately forecast the future sea surface temperature. There are other factors that impact the sea surface temperature. Thus, we added the exogenous variable to increase the prediction accuracy. The exogenous variable in this case is the temperature at 10m-90m below the sea surface.

```
> # adding addition variables without feature selection
> train_additional=ts(ndata[c(2:10)], start = c(2004,1,1), end = c(2015,5,15), frequency = 11)
> test_additional=ts(ndata[c(2:10)], start = c(2015,6,17), end = c(2018,4,12), frequency = 11)
> |
```

Then, we utilized the `auto.arima` function to pick the best model. The best model turned out to be `arima(0,1,0)`.

```
> fitauto_seasonal1<-auto.arima(train_0m,xreg=train_additional, seasonal=TRUE,trace=TRUE,test='kpss',ic='bic')
```

```
Regression with ARIMA(2,1,2)(1,0,1)[11] errors : Inf
Regression with ARIMA(0,1,0) errors : 4.051812
Regression with ARIMA(1,1,0)(1,0,0)[11] errors : 8.392856
Regression with ARIMA(0,1,1)(0,0,1)[11] errors : 6.419628
ARIMA(0,1,0) : -0.7662402
Regression with ARIMA(0,1,0)(1,0,0)[11] errors : 7.529331
Regression with ARIMA(0,1,0)(0,0,1)[11] errors : 7.002356
Regression with ARIMA(0,1,0)(1,0,1)[11] errors : 10.3817
Regression with ARIMA(1,1,0) errors : 5.697884
Regression with ARIMA(0,1,1) errors : 5.236153
Regression with ARIMA(1,1,1) errors : Inf
```

```
Best model: Regression with ARIMA(0,1,0) errors
```

With the best model picked, we used the `forecast` function to predict the sea surface temperature of next 32 months. Below is how the result looks.

```
> forecastedvalue1=forecast(fitauto_seasonal1,xreg=test_additional,h=32)
> forecastedvalue1
```

	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2015.455		17.38864	17.12408	17.65319	16.98403	17.79324
2015.545		17.31676	16.94262	17.69090	16.74457	17.88896
2015.636		17.43723	16.97900	17.89545	16.73643	18.13802
2015.727		17.41792	16.88880	17.94703	16.60871	18.22712
2015.818		18.79294	18.20138	19.38451	17.88822	19.69766
2015.909		21.71286	21.06484	22.36089	20.72179	22.70393
2016.000		23.49975	22.79980	24.19970	22.42927	24.57023
2016.091		23.60324	22.85496	24.35152	22.45885	24.74763
2016.182		22.65849	21.86482	23.45215	21.44468	23.87230
2016.273		20.64776	19.81116	21.48436	19.36829	21.92723
2016.364		18.17316	17.29573	19.05060	16.83125	19.51508
2016.455		16.90580	15.98935	17.82225	15.50422	18.30739
2016.545		16.16461	15.21074	17.11848	14.70579	17.62343
2016.636		16.36347	15.37360	17.35335	14.84959	17.87736
2016.727		17.32547	16.30085	18.35009	15.75844	18.89249
2016.818		18.95617	17.89794	20.01439	17.33775	20.57458
2016.909		20.90323	19.81244	21.99402	19.23501	22.57145
2017.000		22.53833	21.41591	23.66074	20.82174	24.25491
2017.091		22.72026	21.56709	23.87343	20.95663	24.48388
2017.182		21.25678	20.07365	22.43991	19.44734	23.06622
2017.273		20.11404	18.90169	21.32638	18.25991	21.96816
2017.364		17.74546	16.50458	18.98634	15.84770	19.64322
2017.455		16.57032	15.30155	17.83909	14.62991	18.51073
2017.545		15.69027	14.39421	16.98632	13.70812	17.67241
2017.636		16.14639	14.82361	17.46917	14.12338	18.16941
2017.727		17.19401	15.84504	18.54299	15.13093	19.25709
2017.818		19.40085	18.02617	20.77552	17.29847	21.50323
2017.909		20.83899	19.43910	22.23889	18.69803	22.97995
2018.000		22.50521	21.08053	23.92988	20.32635	24.68406
2018.091		22.85289	21.40386	24.30193	20.63679	25.06900
2018.182		22.41432	20.94133	23.88730	20.16158	24.66706
2018.273		21.02429	19.52773	22.52084	18.73551	23.31307

The next step is to assess prediction accuracy. The MSPE is 0.3109 for ARIMAX, big improvement compared to initial model without exogenous variable

```
> mean((forecastedvalue1$mean-test_0m)^2)
[1] 0.3108831
```

Furthermore, we wanted to consider feature selection to see if we can improve prediction accuracy. We've run the linear regression, according to summary output, p-values of 10m, 80m, and 90m are low compared to other features.

```
> # adding addition variables with feature selection
> train<-ndata[1:126,]
> test<-ndata[127:158,]
> attach(train)
```

```
> linearreg =lm("0m"~., data=train)
>
> summary(linearreg)

Call:
lm(formula = "0m" ~ ., data = train)

Residuals:
    Min       1Q   Median       3Q      Max
-0.93251 -0.08856  0.00844  0.13463  0.46672

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -1.24818    0.58636  -2.129  0.03539 *
      10m      1.36090    0.17045   7.984 1.14e-12 ***
      20m     -0.17439    0.30943  -0.564  0.57412
      30m     -0.25180    0.25254  -0.997  0.32081
      40m      0.06504    0.16504   0.394  0.69426
      50m      0.02976    0.19487   0.153  0.87887
      60m     -0.12451    0.28781  -0.433  0.66609
      70m     -0.29390    0.25331  -1.160  0.24833
      80m     -0.64705    0.23798  -2.508  0.01352 *
      90m      0.52371    0.16448   3.184  0.00187 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2411 on 116 degrees of freedom
Multiple R-squared:  0.9899,    Adjusted R-squared:  0.9891
F-statistic: 1262 on 9 and 116 DF, p-value: < 2.2e-16
```

With selected features of 10m, 80m, and 90m, we reevaluated our model using auto.arima function, the best model is still arima(0,1,0)

```
> train_FeatureSelection<-ts(ndata[,c(2,9,10)], start = c(2004,1,1), end = c(2015,5,15), frequency = 11)
> test_FeatureSelection=ts(ndata[,c(2,9,10)], start = c(2015,6,17), end = c(2018,4,12), frequency = 11)
>
> fitauto_seasonal2<-auto.arima(train_0m,xreg=train_FeatureSelection, seasonal=TRUE,trace=TRUE,test='kpss',ic='bic')

Regression with ARIMA(2,1,2)(1,0,1)[11] errors : Inf
Regression with ARIMA(0,1,0) errors : -3.999919
Regression with ARIMA(1,1,0)(1,0,0)[11] errors : 3.532423
Regression with ARIMA(0,1,1)(0,0,1)[11] errors : 3.181162
ARIMA(0,1,0) : -8.823815
Regression with ARIMA(0,1,0)(1,0,0)[11] errors : 0.1336801
Regression with ARIMA(0,1,0)(0,0,1)[11] errors : 0.1310193
Regression with ARIMA(0,1,0)(1,0,1)[11] errors : 4.95854
Regression with ARIMA(1,1,0) errors : -0.211681
Regression with ARIMA(0,1,1) errors : -0.4297958
Regression with ARIMA(1,1,1) errors : Inf

Best model: Regression with ARIMA(0,1,0) errors
```

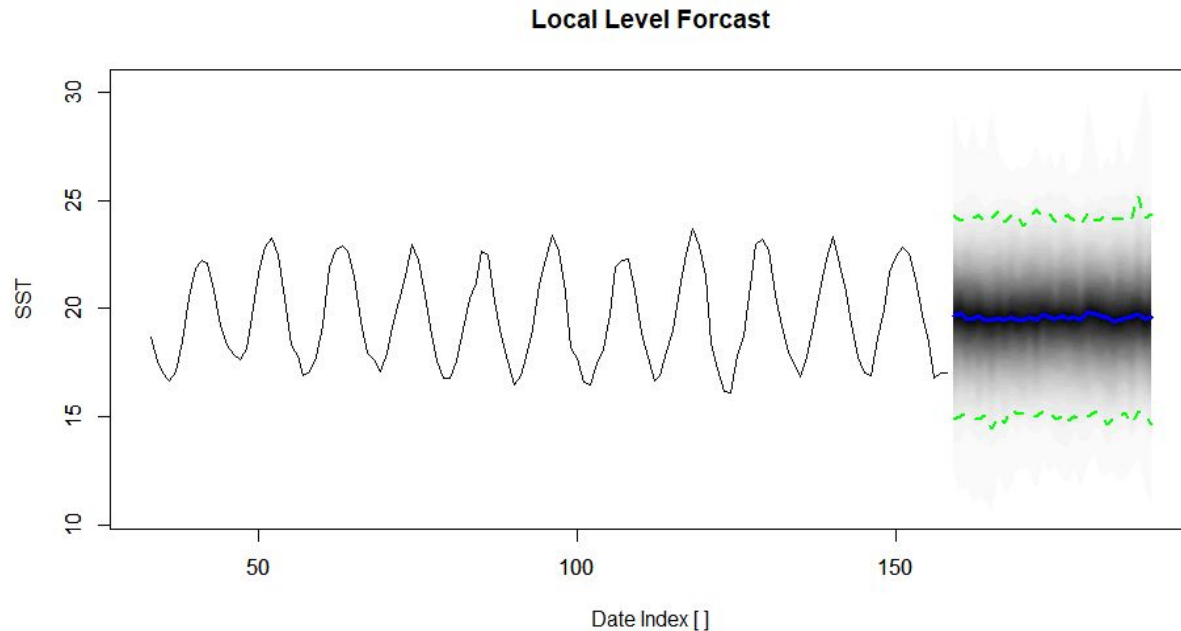
Again, we predicted sea surface temperature using forecast function, and assessed the MSPE. MSPE score turned out to be 0.2894. A slight improvement compared to the model without feature selection.

```
> forecastedvalue2=forecast(fitauto_seasonal2,xreg=test_FeatureSelection,h=32)
> mean((forecastedvalue2$mean-test_0m)^2)
[1] 0.2893556
>
```

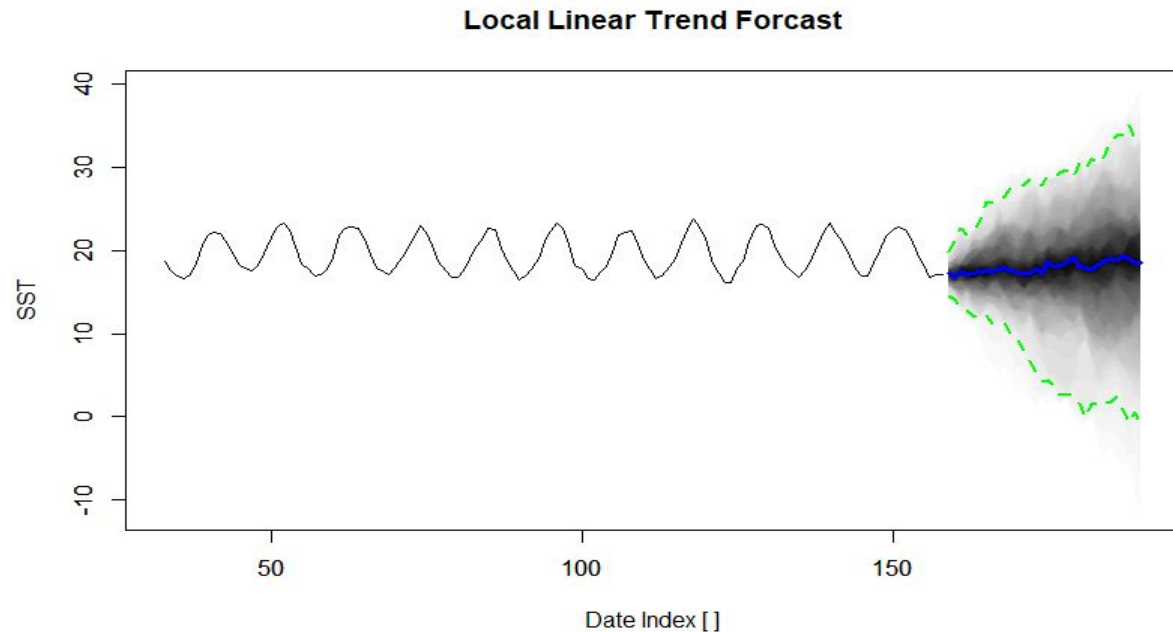
Model 3: Bayesian Structural Time Series (BSTS) (Chris)

BSTS is a time series regression method that uses dynamic linear models fit using Markov Chain Monte Carlo. With our data, we first started off with the most simple useful model in a structured time series and then worked our way up with increasing the complexity of the models. In the BSTS Models, the dark blue line is the median of the predictive distributions, the green lines indicate the prediction intervals (default is the 95% prediction interval), and the grey shaded area represents the posterior density distributions.

The most simple useful model in structured time series is the Local Level. BSTS Model of the Local Level forecasts is based around the average value of recent observations. The Local Level is also a random walk observed in noise.

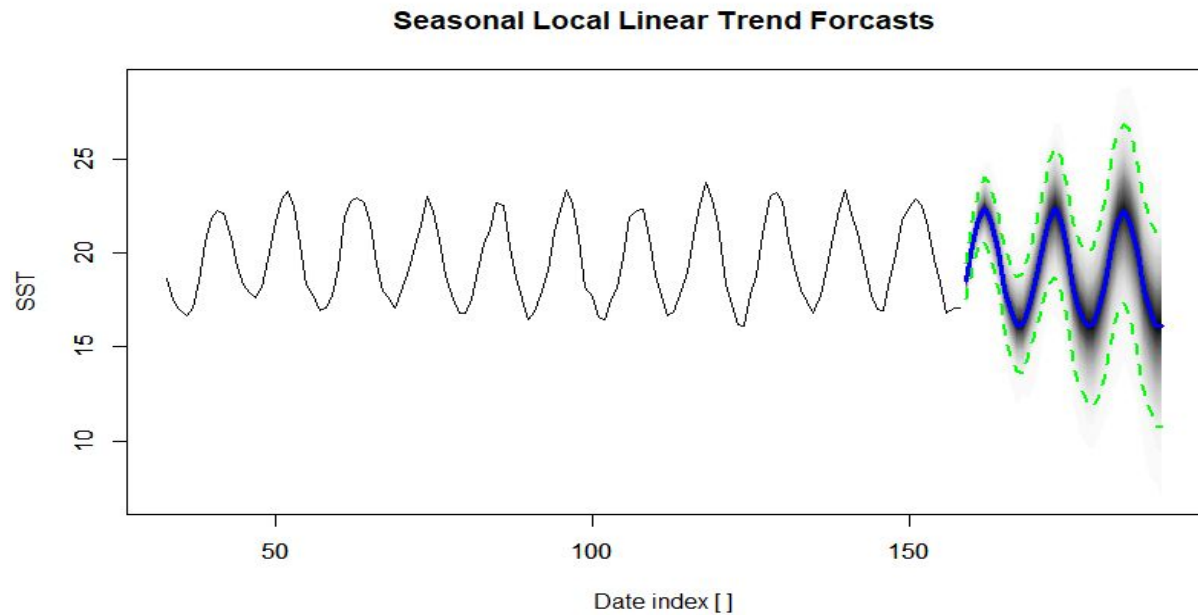


With the local level, we got a MSPE of 5.09. After attempting the local level model, we increased the complexity by adding a local linear trend component.




In our attempt of adding a local linear trend component, our MSPE score got worse, which makes sense. It was a MSPE of 9.955. As you can see in our data, we have an obvious seasonality. The Local Linear Trend model attempts to capture the trend at the tail ends of the time series. It is especially ideal for time series that have a particular direction and if we wanted the forecast to reflect on the past observations.

After adding on the local linear component, we then again increased our model's complexity by adding a seasonal component. With adding a seasonal component, we expect our model to fit pretty well.




As we can see, after adding the seasonal component to make our forecasts seasonal, our new model is now much better than the previous two models since the two prior ones failed to characterize seasonality. We now have a MSPE of 2.25.

Now, after having a good model for our data, we wanted to attempt to make our model even better, and to do so, we created a regressional component by setting SST/0m, as our output, and the underlying layers (10m - 90m) as our features. We also look at the trend components of our data to let us know what may or may not be a good expected model size.



Looking at the trend components, the white bars represent positive betas, and the black bars represent negative betas for our regressional component. We see that 10m and 20m both are positive and have a high inclusion probability, so we expect that an expected model size of 1 and or 2 will be optimal. In our case, we set our expected model size to 1. The expected model size sets our spike and slab prior to having one spike, in other words, we are expecting our one depth to heavily influence the sea temperature, most likely the layer directly underneath (10m).



As we can see, adding a regressional component narrowed our prediction intervals improving our model. With the narrowed prediction intervals and improved performance, our new MSPE is now 0.024.

Conclusion (Yege)

The aim of the time series model is to predict the future values for the series by studying the past observations. In this project, we fitted the following three models, SARIMA model without including additional variable, SARIMA model with additional variable and BSTS. The SARIMA model performed much better after we added the additional variables. However, we only see a slight improvement when feature selection was applied. Among the most effective approaches for analyzing time series data, BSTS seasonal local linear trend forecast with regression is employed with the lowest MSPE and appropriate model was adaptively formed based on the given data.